

CS309 ASSIGNMENT1

Krishanu Saini - (190001029)

September 11, 2021

1 Problem Statement

In this assignment, you will develop MPI program for the following tasks.

1. Write an MPI program with 2/4 processes for multiplying two vectors of size 100k (5 marks).
2. Write an MPI program for multiplying two random matrices of size 5k X 5k (5 marks).

You need to print runtime of your code as well.

2 Multiplying 2 vectors

Suppose we have 2 vectors $A_{n \times 1}$ and $B_{n \times 1}$. We know product of 2 vectors is $A^T B = A_1 \cdot B_1 + A_2 \cdot B_2 + A_3 \cdot B_3 + A_4 \cdot B_4$.

To make it parallel using mpi we will use **data parallelism** steps below.

1. Calculate size of each block = $n / \text{number of process}$.
2. Split A and B into array of size = size_block
3. Accumulate the result in master process using MPI_reduce.

Time Complexity: $O(n/p)$.

3 Code

1. Multiplying 2 vectors : Link

```
#include<stdio.h>
#include<stdlib.h>
#include<mpi.h>

int main(){
    MPI_Init(NULL, NULL);
    int world_rank, world_size;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // input 100k size vector
    int n = 100000;
    int *A = (int*)malloc(n*sizeof(int));
    int *B = (int*)malloc(n*sizeof(int));

    // fill data in arr
    for(int i=0;i<n;i++){
        A[i] = (i % 100);
        B[i] = (i % 50);
    }

    int size = n / world_size;

    int *subarray_A = (int*)malloc(size*(sizeof(int)));
    int *subarray_B = (int*)malloc(size*(sizeof(int)));

    MPI_Scatter(A, size, MPI_INT, subarray_A, size, MPI_INT, 0,
        MPI_COMM_WORLD);
    MPI_Scatter(B, size, MPI_INT, subarray_B, size, MPI_INT, 0,
        MPI_COMM_WORLD);

    // dot product here
    int *resultBuf = NULL;
    if(world_rank == 0){
        int *resultBuf = (int*)malloc(world_size*(sizeof(int)));
    }
}
```

```
int sub_sum = 0;
for(int i=0;i<size;i++){
    sub_sum += subarray_A[i]*subarray_B[i];
}

printf("Processor %d gives sum %d\n", world_rank, sub_sum);
MPI_Barrier(MPI_COMM_WORLD);

int ans = 0;
MPI_Reduce(&sub_sum, &ans, 1, MPI_INT, MPI_SUM, 0,
    MPI_COMM_WORLD);

if(world_rank == 0){
    // gather here
    printf("Vector multiplication gives %d\n", ans);
}

MPI_Finalize();
return 0;
}
```

4 OUTPUT

Explanation - For vectors $A = [0, 1, 2, 3]^T$ and $B = [0, 1, 2, 3]^T$
the vector multiplication = $A^T \cdot B$ ie $z = 0 \cdot 0 + 1 \cdot 1 + 2 \cdot 2 + 3 \cdot 3 = 0 + 1 + 4 + 9 = 14$

2 X Processes

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$ mpicc vector_mult.c -o vector_mult
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$ mpirun -n 2 ./vector_mult
0 1 2 3
0 1 2 3
Processor 0 gives sum 1
Processor 1 gives sum 13
Vector multiplication gives 14
```

4 X Processes

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$ mpirun -n 4 ./vector_mult
0 1 2 3
0 1 2 3
Processor 0 gives sum 0
Processor 2 gives sum 4
Processor 3 gives sum 9
Processor 1 gives sum 1
Vector multiplication gives 14
```

Now for 100K sized vector.

This is for input specified in code. $A = [i : i\%100]^T$, $B = [i : i\%50]^T$.

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$ mpirun -n 2 ./vector_mult
Processor 0 gives sum 71050000
Processor 1 gives sum 71050000
Vector multiplication gives 142100000
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$ ls
```

5 Multiplying 2 matrices

Suppose we have 2 square matrices $A_{n \times n}$ and $B_{n \times n}$. We know product of 2 matrices is $O(n^3)$ Operation using 3 for loops.

To make it parallel using mpi we will use **data parallelism** steps below.

1. Split $A_{n \times n}$ into multiple parts having row size $r = n / \text{number of processors}$.
as $A[1...r][n]$, $A[r + 1...2r][n]$ so on.
2. MPI_Send $A_{r \times n}$ and B matrix to multiple processes.
3. Calculate matrix $C_{r \times n} = A_{r \times n} \cdot B_{n \times n}$ using matrix multiplication.
4. Master process gathers all $C_{r \times n}$'s in final matrix C.

Time Complexity: $O(n^3/p)$.

6 Code

2. Multiplying 2 matrices : [Link](#)

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define MATSIZE 4
#define NRA MATSIZE
#define NCA MATSIZE
#define NCB MATSIZE
#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2

int main(int argc, char *argv[])
{
    int numtasks,
        taskid,
        numworkers,
```

```

    source,
    dest,
    mtype,
    rows,
    averow, extra, offset,
    i, j, k, rc;
double a[NRA][NCA],
    b[NCA][NCB],
    c[NRA][NCB];
MPI_Status status;

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &taskid);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
if (numtasks < 2)
{
    printf("Need at least two MPI tasks. Quitting...\n");
    MPI_Abort(MPI_COMM_WORLD, rc);
    exit(1);
}
numworkers = numtasks - 1;
if (taskid == MASTER)
{
    printf("mpi_mm has started with %d tasks.\n", numtasks);
    printf("Initializing arrays...\n");
    for (i = 0; i < NRA; i++)
        for (j = 0; j < NCA; j++)
            a[i][j] = i + j;
    for (i = 0; i < NCA; i++)
        for (j = 0; j < NCB; j++)
            b[i][j] = i * j;
    double start = MPI_Wtime();
    averow = NRA / numworkers;
    extra = NRA % numworkers;
    offset = 0;
    mtype = FROM_MASTER;
    for (dest = 1; dest <= numworkers; dest++)
    {
        rows = (dest <= extra) ? averow + 1 : averow;
        printf("Sending %d rows to task %d offset=%d\n", rows, dest,

```

```

        offset);
    MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
    MPI_Send(&a[offset][0], rows * NCA, MPI_DOUBLE, dest, mtype,
             MPI_COMM_WORLD);
    MPI_Send(&b, NCA * NCB, MPI_DOUBLE, dest, mtype,
             MPI_COMM_WORLD);
    offset = offset + rows;
}
mtype = FROM_WORKER;
for (i = 1; i <= numworkers; i++)
{
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
             &status);
    MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD,
             &status);
    MPI_Recv(&c[offset][0], rows * NCB, MPI_DOUBLE, source, mtype,
             MPI_COMM_WORLD, &status);
    printf("Received results from task %d\n", source);
}
printf("Result Matrix:\n");
for (i = 0; i < NRA; i++)
{
    printf("\n");
    for (j = 0; j < NCB; j++)
        printf("%6.2f ", c[i][j]);
}
double finish = MPI_Wtime();
printf("Done in %f seconds.\n", finish - start);
}
if (taskid > MASTER)
{
    mtype = FROM_MASTER;
    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
             &status);
    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,
             &status);
    MPI_Recv(&a, rows * NCA, MPI_DOUBLE, MASTER, mtype,
             MPI_COMM_WORLD, &status);
}

```

```

MPI_Recv(&b, NCA * NCB, MPI_DOUBLE, MASTER, mtype,
        MPI_COMM_WORLD, &status);

for (k = 0; k < NCB; k++)
    for (i = 0; i < rows; i++)
    {
        c[i][k] = 0.0;
        for (j = 0; j < NCA; j++)
            c[i][k] = c[i][k] + a[i][j] * b[j][k];
    }
mtype = FROM_WORKER;
MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
MPI_Send(&c, rows * NCB, MPI_DOUBLE, MASTER, mtype,
        MPI_COMM_WORLD);
}
MPI_Finalize();
}

```

7 OUTPUT

Matrix A =

0	1	2	3
1	2	3	4
2	3	4	5
3	4	5	6

Matrix B =

0	0	0	0
0	1	2	3
0	2	4	6
0	3	6	9

```

PROBLEMS 6 OUTPUT DEBUG CONSOLE TERMINAL
Initializing arrays...
Sending 2 rows to task 1 offset=0
Sending 1 rows to task 2 offset=2
Sending 1 rows to task 3 offset=3
Received results from task 1
Received results from task 2
Received results from task 3
Result Matrix:

0.00 14.00 28.00 42.00
0.00 20.00 40.00 60.00
0.00 26.00 52.00 78.00
0.00 32.00 64.00 96.00 Done in 0.000106 seconds.
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$

```

Matrix A =

0	1	2
1	2	3
2	3	4

Matrix B =

0	1	2
1	2	3
2	3	4

```

mpi mm has started with 4 tasks.
Initializing arrays...
Sending 1 rows to task 1 offset=0
Sending 1 rows to task 2 offset=1
Sending 1 rows to task 3 offset=2
Received results from task 1
Received results from task 2
Received results from task 3
Result Matrix:

5.00 8.00 11.00
8.00 14.00 20.00
11.00 20.00 29.00 Done in 0.000059 seconds.
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab2$

```