

# CS309 ASSIGNMENT2

Krishanu Saini - (190001029)

September 17, 2021

## 1 Problem Statement

In this assignment, you will develop MPI program for the following tasks.

1. Parallel merge sort starts with  $n/\text{comm\_size}$  keys assigned to each process. It ends with all the keys stored on process 0 in sorted order. To achieve this, it uses the same tree-structured communication that we used to implement a global sum.

However, when a process receives another process' keys, it merges the new keys into its already sorted list of keys.

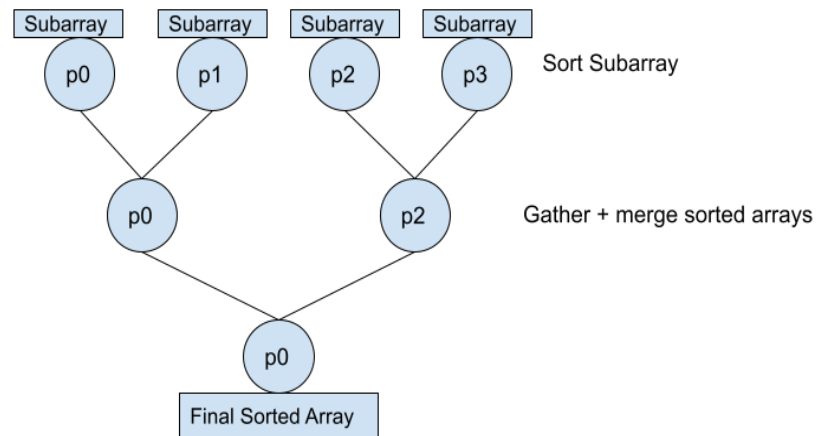
Notes

- (a) Write a program that implements parallel mergesort. Process 0 should read in  $n$  and broadcast it to the other processes.
- (b) Each process should use a random number generator to create a local list of  $n/\text{comm\_size}$  ints.
- (c) Each process should then sort its local list, and process 0 should gather and print the local lists.
- (d) Then the processes should use tree-structured communication to merge the global list onto process 0, which prints the result. (MPI, 20 marks)

## 2 Solution

### Algorithm

1. Input size of array -  $n$ . Broadcast it
2. Generate random array of  $n/world\_size$  integers
3. Sort subarray in each processor.
4. Merging to be done like balanced binary tree
5. Print result in master processor (0).



## 3 Analysis

### 3.1 Time Complexity

Let  $n$  be number of elements,  $p$  be the number of processors.

1. Input of array  $O(n/p)$  per processor
2. Sort subarray  $O(n/p \cdot \log(n/p))$
3. Merging sorted arrays  $O(n)$   
There are  $\log_2(p)$  levels. Each level requires  $O(2 \cdot \text{size}(\text{prev\_subarray}))$ .  
ie  $T = 2n/p \cdot (1 + 2 + \dots + 2^{\log(p)-1}) = 2n/p \cdot 2^{\log(p)} - 1 = 2n$
4. Communication time =  $\log(p)$

Net Time Complexity =  $O(n/p + \log(p) + n/p \cdot \log(n))$

### 3.2 Space Complexity

1. Input array total  $O(n)$
2. Subarray buffer  $O(n/\text{world\_size})$
3. Temporary arrays max  $n/p \cdot (p + p/2 + p/4 + \dots + \log(p) \text{ times}) = O(n)$

Net Space Complexity =  $O(n)$

## 4 Code

---

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <mpi.h>

int randInt()
{
    return rand() % 1000;
}

/***** Merge Function *****/
void merge(int *a, int *b, int l, int m, int r)
{
    // printf("[ ] [ ] array - ");
    // for(int i=l;i<=r;i++){
    //     printf("%d ", a[i]);
    // }
    // printf("\n");
    int h, i, j, k;
    h = l;
    i = l;
    j = m + 1;

    while ((h <= m) && (j <= r))
    {
        if (a[h] <= a[j])
        {
            b[i] = a[h];
            h++;
        }

        else
        {

```

```

        b[i] = a[j];
        j++;
    }

    i++;
}

if (m < h)
{

    for (k = j; k <= r; k++)
    {

        b[i] = a[k];
        i++;
    }
}

else
{

    for (k = h; k <= m; k++)
    {

        b[i] = a[k];
        i++;
    }
}

for (k = l; k <= r; k++)
{

    a[k] = b[k];
}
}

/***** Recursive Merge Function *****/
void mergeSort(int *a, int *b, int l, int r)
{

```

```

int m;

if (l < r)
{
    m = (l + r) / 2;

    mergeSort(a, b, l, m);
    mergeSort(a, b, (m + 1), r);
    merge(a, b, l, m, r);
}
}

int main(int argc, char *argv[])
{
    int world_rank, world_size;
    double time1, time2,duration;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    srand(world_rank+1);
    /*----- Take Input -----*/
    int n;
    if (world_rank == 0)
    {
        printf("Input array size(n): ");
        scanf("%d", &n);
    }

    time1 = MPI_Wtime();
    /*----- Broadcast -----*/
    MPI_Bcast(&n, sizeof(int), MPI_INT, 0, MPI_COMM_WORLD);

    /*----- Scatter -----*/
    if (n % world_size != 0)
    {
        printf("Use correct number of processors\n");
        exit(0);
    }
}

```

```

}

int r = n / world_size;
int *subarray = (int *)calloc(r, sizeof(int));
int *sorted = (int *)calloc(r, sizeof(int));

int sizeSorted = r;

// MPI_Scatter(arr, r, MPI_INT, subarray, r, MPI_INT, 0,
    MPI_COMM_WORLD);

printf("Processor %d: subarray: ", world_rank);
for (int i = 0; i < r; i++)
{
    subarray[i] = randInt();
    printf("%d ", subarray[i]);
    sorted[i] = subarray[i];
}
printf("\n");

mergeSort(subarray, sorted, 0, r-1);

/*----- tree based merging -----*/
/*
    n n n n n n n n
    | / | / | / | /
    n  n  n  n
    |  /   |  /
    n      n
    | ____/
    n
*/

int cur_world_size = world_size;
int np = 2;
while(cur_world_size > 1){
    if(world_rank % (np/2) != 0) {
        break;
    }
    int r;

```

```

int *other_sorted = (int* )malloc(sizeSorted*sizeof(int));
if(world_rank % np == (np/2)){
    r = sizeSorted;
    for(int i=0;i<r;i++){
        other_sorted[i] = sorted[i];
    }
    MPI_Send(&r, 1, MPI_INT, world_rank-np/2, 5, MPI_COMM_WORLD);
    MPI_Send(other_sorted, r, MPI_INT, world_rank-np/2, 11,
        MPI_COMM_WORLD);
} else if(world_rank/(np/2) != cur_world_size-1){ // last one
    excluded if cur_world_size odd
    MPI_Recv(&r, 1, MPI_INT, world_rank+np/2, 5, MPI_COMM_WORLD,
        MPI_STATUS_IGNORE);
    MPI_Recv(other_sorted, r, MPI_INT, world_rank+np/2, 11,
        MPI_COMM_WORLD, MPI_STATUS_IGNORE);

    /*----- merging 2 sorted arrays -----*/
    int r1 = sizeSorted;
    sorted = (int* )realloc(sorted, (r+r1)*sizeof(int));
    sizeSorted = r+r1;
    for(int i=0;i<r;i++){
        sorted[r1+i] = other_sorted[i];
    }
    int* newSorted = (int*)malloc(sizeSorted*sizeof(int));
    merge(sorted, newSorted, 0, r1-1, r+r1-1);
    sorted = newSorted;

}
// MPI_Barrier(MPI_COMM_WORLD);
np *= 2;
cur_world_size = (cur_world_size+1) / 2;
}

if(world_rank == 0){
    time2 = MPI_Wtime();

    duration = time2 - time1;

    printf("Sorted array - ");
    int r = sizeSorted;

```



```
    for(int i=0;i<r;i++){
        printf("%d ", sorted[i]);
    }
    printf("\n");
    printf("Time taken by program = %0.9f\n", duration*1e6);

}

MPI_Barrier(MPI_COMM_WORLD);
MPI_Finalize();
return 0;
}
```

---

## 5 OUTPUT

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab3$ mpicc merge_sort_mpi.c -o merge_sort_mpi
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab3$ mpirun -n 2 ./merge_sort_mpi
Input array size(n): 4
Array: 560 568 921 968
Processor 0: subarray: 560 568
Processor 1: subarray: 921 968
Sorted array - 560 568 921 968
```

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab3$ mpirun -n 4 ./merge_sort_mpi
Input array size(n): 8
Array: 322 275 193 599 79 336 97 43
Processor 0: subarray: 322 275
Processor 1: subarray: 193 599
Processor 2: subarray: 79 336
Processor 3: subarray: 97 43
Sorted array - 43 79 97 193 275 322 336 599
```

```
krishanu2001@LAPTOP-V4CKFTKN:/mnt/c/Users/krishanu/Desktop/sem5/PARALLEL/mpi/lab3$ mpirun -n 4 ./merge_sort_mpi
Input array size(n): 16
Array: 239 963 707 228 52 217 514 875 820 177 491 210 375 761 678 51
Processor 0: subarray: 239 963 707 228
Processor 1: subarray: 52 217 514 875
Processor 2: subarray: 820 177 491 210
Processor 3: subarray: 375 761 678 51
Sorted array - 51 52 177 210 217 228 239 375 491 514 678 707 761 820 875 963
```

## 6 Explanation

Comparing parallel algorithm runtime vs serial.

n	p	Parallel( $\mu s$ )	Serial( $\mu s$ )
10	2	32	45
100	4	40	100
1k	4	87	144
10k	4	1029	3250
100k	4	4646	11544

Note: We observe speedup is not exactly p times. There is some overhead due to communication and memory delays. For larger problem size - speedup is better.

Comparing Speedup vs processors

n	100	1k	10k	100k
p=1	14( $\mu s$ )	98	1133	13025
p=2	35	93	1030	6945
p=4	65	83	472	4716
p=10	95	105	415	4344

S vs P

